

CS2102 MiniProject-2026

IITPSx Processor

The aim of this mini project is to write a two pass assembler for the IITPSx processor instruction set. Then write and test programs in IITPSx assembly. A final part is to write an emulator for the IITPSx processor (replacing the one provided).

Tasks

1. Write a two pass assembler for the assembly language.
The assembler must,
 - Read assembly language from a text file, assigning label values and instruction opcodes. The format of the assembly language is described.
 - Diagnose common assembly errors such as unknown instruction, no such label, duplicate label.
 - Produce an object file of the produced machine code. This file should be a binary file. Code starts at address zero.
 - Produce a listing file. There is a choice of the format of the listing file. It can either be a simple memory dump, or show the bytes produced for each instruction, and that instruction's mnemonic. The formats are shown.
 - You must write the assembler in ISO C89. ('gcc -std=c89 -pedantic -W -Wall' is a good way to check things.) or others are also fine.
2. Test your assembler with the sample programs listed.
3. Test your assembler with additional programs and submit evidence of this.
4. Write a bubble sort program in IITPSx assembly. The start of this file is provided, you have to fill in the blanks.
5. Write an emulator for the IITPSx processor. This should have some, but not need all, of the functionality of the `emu` program provided.
6. Implement and test the two new IITPSx arithmetic instructions, `mul` and `div`, in both the assembler and the emulator.

Sample `asm` and `emu` (executable) are given.

IITPSx Assembly Language

This assembly language is for the IITPSx processor, which has four registers,

- Two registers, A & B, arranged as an internal **stack**.
- A program counter, PC
- A stack pointer, SP

These registers are 32 bits in size. Instructions have either no operands or a single operand. The operand is a signed 2's complement value. The encoding uses the bottom 8 bits for opcode and the upper 24 bits for operand.

As with most assembly languages, this is line based (one statement per line). Comments begin with a ';' and anything on the line after the ';' is ignored. Blank lines and lines containing only a comment are permitted (and ignored). White space (' and tabs) are permitted at the beginning of a line (and ignored). Label definitions consist of the label name followed by a ':', and an optional statement (there is not necessarily a space between the ':' and the statement). A label use is just the label name. For branch instructions label use should calculate the branch displacement. For non-branch instructions, the label value should be used directly. A valid label name is an alphanumeric string beginning with a letter . An operand is either a label or a number, the number can be decimal, hex or octal.

The following are all permitted lines

```
; a comment
                ; another comment
label1:         ; a label on its own
ldc 5           ; an instruction
label2: ldc 5   ; a label and an instruction
                adc 5 ; an instruction
label3:ldc label3 ;look no space between label and mnemonic
```

Each statement consists of a mnemonic (instruction name) and an optional operand (number or label).

IITPSx Instruction Set

The instruction semantics do not show the incrementing of the PC to the next instruction. This is implicitly performed by each instruction *before* the actions of the instruction are done. The IITPSx processor adds two required arithmetic instructions, mul and div, using opcodes 19 and 20.

Mnemonic	Opcode	Operand	Formal Specification	Description
data		value		Reserve a memory location, initialized to the value specified
ldc	0	value	B := A; A := value;	Load accumulator with the value specified
adc	1	value	A := A + value;	Add the value specified to the accumulator

ldl	2	offset	B := A; A := memory[SP + offset];	Load local
stl	3	offset	memory[SP + offset] := A; A := B;	Store local
ldnl	4	offset	A := memory[A + offset];	Load non-local
stnl	5	offset	memory[A + offset] := B;	Store non-local
add	6		A := B + A;	Addition
sub	7		A := B - A;	Subtraction
shl	8		A := B << A;	Shift left
shr	9		A := B >> A;	Shift right
adj	10	value	SP := SP + value;	Adjust SP
a2sp	11		SP := A; A := B	Transfer A to SP;
sp2a	12		B := A; A := SP;	Transfer SP to A
call	13	offset	B := A; A := PC; PC := PC + offset;	Call procedure
return	14		PC := A; A := B;	Return from procedure
brz	15	offset	if A == 0 then PC := PC + offset;	If accumulator is zero, branch to specified offset
brlz	16	offset	if A < 0 then PC := PC + offset;	If accumulator is less than zero, branch to specified offset
br	17	offset	PC := PC + offset;	Branch to specified offset
HALT	18			Stop the emulator. This is not a 'real' instruction, but needed to tell your emulator when to finish.
mul	19		A := B * A;	Multiplication
div	20		if A == 0 then runtime error; else A := B / A;	Signed integer division (B divided by A); division by zero must be reported by the emulator
SET		value		Set the label on this line to the specified value (rather than the PC). This is an optional extension, for which additional marks are available.

Listing File Format

The listing file is produced by the assembler and is a human readable file showing what value is stored at each address. The format is an address followed by zero or one 32 bit values (as 8 hex characters).

With the output, you can choose to show the human readable mnemonic and operand, that each instruction corresponds to. You can also show labels, by simply listing the address followed by no data bytes.

Here are some acceptable example outputs,

```
00000000 00000111
00000001 00005AB4
00000002 00006500
00000003 00009D01
```

Or, with 4 locations per line,

```
00000000 00000111 00005AB4 00006500 00009D01
```

When showing labels and mnemonics, the output could be (this is the preferred way of showing things)

```
00000000 00000111 br start
00000001 00005AB4 data 0x5ab4
00000002          start:
00000002 00006500 ldc 0x65
00000003 00009D01 adc 0x9d
```

Example Programs

This is a valid, but nonsense assembly file. Your assembler should not issue any errors (it could issue warning though).

```
; test1.asm
label: ; an unused label
        ldc 0
        ldc -5
        ldc +5
loop: br loop ; an infinite loop
br next ;offset should be zero
next:
        ldc loop ; load code address
        ldc var1 ; forward ref
var1: data 0 ; a variable
```

It is shown as text below.

This example contains many errors. Your assembler should spot them all (it need not copy the error message exactly, but should issue something appropriate).

```

; test2.asm
; Test error handling
label:
label: ; duplicate label definition
br nonesuch ; no such label
ldc 08ge ; not a number
ldc ; missing operand
add 5 ; unexpected operand
ldc 5, 6; extra on end of line
0def: ; bogus label name
fibble; bogus mnemonic
0def ; bogus mnemonic

```

If you implement the SET pseudo instruction, this program should assemble

```

; test3.asm
; Test SET
val: SET 75
ldc    val
adc    val2
val2: SET 66

```

The following program demonstrates the two IITPSx arithmetic extensions. It should assemble and, when executed, leave 21 in register A.

```

; test4.asm
; Test IITPSx mul and div
ldc 6
ldc 7
mul      ; A = 42
ldc 2
div      ; A = 21
HALT

```

Notes

Binary file

Files may be either text or binary. The default is text. The manual page for `fopen` explains how to open a file in binary mode.

Two Passes

A two pass assembler naturally needs to scan the source file twice. You can and should use a single routine to do both passes, provided that on the first pass it outputs no code, and doesn't fail on undefined labels, and that on the second pass it outputs code and does fail on an undefined label. This is a better solution than having two different routines for each pass (because of code reuse).

One Read

Alternatively you could read the program into an internal form, and then process that twice. This uses less file IO, at the cost of using more memory. If you do this, the internal form should be more advanced than simply storing each line of the file as a string. It is a more complicated solution.

Reading a Line

The functions `gets` and `fscanf` should be handled with care. `gets` suffers from buffer overrun -- if the lines are longer than the programmer expected, it will merrily trash subsequent locations. `fscanf` can also suffer the same problem with reading strings. However, it has a worse flaw, in that it does not recover well from errors. If the line does not contain what `fscanf`'s format string expects, it is very hard to resynchronize. You should use something like `fgets` to read a buffer and then `sscanf` or `strtok`, for a robust program.

Also `getchar` and its ilk return an `int`, not a `char`. It needs to do this because it needs to return all the possible `char` values, plus another value to mean 'end of file'. So storing the return value into a `char` variable is **wrong**.

Tables

If you define your instructions in a suitable data structure, you can define all the instructions in a table. This is better than hard coding particular instruction semantics throughout the program.

Number Format

Numbers in the source file are written as an optional sign, a digit and trailing digits and letters. Not all these will be *valid* numbers. You can use `strtol` to convert a string to a number, and tell you how far it got into the string. It will also figure out whether the number is decimal, hex (begins with `'0x'`), or octal (begins with `'0'`).

Printing Label Values

One of the example listing output shows an instruction printed as `br start`. Firstly, although not mentioned, it would be permissible to show this as `br 75` (if the branch offset is 75), but that wouldn't be as good. As the assembler has replaced label names with their values, how is this achieved? When listing an instruction, the assembler could check if its operand is a PC offset. If so, it could calculate where the branch would go to, and then look in its table of labels to find a label at that address. If it finds one, then use it, otherwise just print the raw offset. Notice, that the assembly source might contain `br 7`. You do not need to check that this actually lands up at an instruction -- if it doesn't it's the assembly language programmer's fault.

ldc result

The `ldc result` instruction in the example program has confused some of you. `ldc` loads a constant, but here I am using it with a label -- why? What I wanted to get hold of was the address of the `result` data location, so I could access it as an array. So, I needed its address, hence the `ldc`. All the assembler does with the operand is lookup the label's value and put that in the instructions operand. *It doesn't understand what an instruction does*. You'll see I access it with `ldnl` instructions later on.

Warnings

What's the difference between errors and warnings? An error is something which is incorrect and prevents completion of the task. A warning is something which is not

strictly incorrect (i.e. is allowed), but is strange or dubious. It does not prevent completion of the task. An example is an unused label. It is not incorrect to declare a label and not use it, however that is a strange thing to do, and might be hiding some more important problem. For instance perhaps the programmer misspelt the label when trying to use it, and accidentally used a different label. This is one reason why variable and function names should be very different to each other. Don't go calling your variables 'var1', 'var2', 'var3' etc!

The Submission

Submission will be electronic, using the following link.

<https://u.pcloud.com/#/puplink?code=F0wXZD4XBYBEvrpVoSK2TR6suKbtX5MAy>

You should submit the following files

- The source code for your assembler. The main (and possibly only) file should be called 'asm.c'.
- The source code for your emulator (if you did that bit). The main (and possibly only) file should be called 'emu.c'.
- Your programs should either be compilable with

```
Gcc asm.c -o asm
gcc emu.c -o emu
```

or, if it is more complicated then submit a `Makefile` which contains the correct rules.

- Test assembly files you used to verify your assembler and emulator. These should be called `test01.asm` and so on.
- Output files from your tests. These should be named `test01.log` and so on.
- The assembler listing files. These should be named `test01.lst` and so on.
- The "claims" text file explained below.

IMPORTANT NOTE:

You will receive **no marks** if you do not deliver the following:

- You **must** submit evidence that your assembler works (its not acceptable to submit only a program listing).
- Also you must submit a text file named "claims" that lists (according to the marking scheme below) the claims you make about your submission (i.e. what do you believe you have done correctly).
- When asked by the marker, you should be able to explain **in deep detail** the inner workings and structure of your code.

The submission deadline is midnight 11th October 2026

The Marking Scheme

General marks

This applies to all the code you submit (i.e. both the assembler and the emulator).

- Correctly submitted source files with your name in them. All files should have your name and user id at the start, along with a **declaration of authorship**. * (and must)
- Compiles without error when being pedantic (`gcc -std=c89 -pedantic`). *
- Compiles without warnings (`gcc -std=c89 -pedantic -W -Wall -Wpointer-arith -Wwrite-strings -Wstrict-prototypes`).
- Consistent and sensible formatting with sensible program structure.
- Sensible variable, function & type names with **explanatory comments**. * (very important issue)

The assembler part (60%)

- Either uses a single process/routine for both passes, or reads into an internal form on first pass. * (you can only do one of these, not both)
- Diagnoses label errors (duplicate, missing, etc). *
- Produces listing file (at least memory dump). *
- Advanced listing file.
- Has a table of instruction names and expected operands. *
- Assembles test program. *
- Additional test programs (for **both** failure and success, include brief explanation as comments within each test file). *
- Implements **and** demonstrates the instruction SET.
- Implements and demonstrates the IITPSx mul and div instructions. *

Bubble sort program,

- Assembles. *
- Works. *

The emulator part (25%)

- Loads object file.
- Produces memory dump.
- Executes test program.
- Executes mul and div correctly, including detection of division by zero. *
- Detects errant programs.

Test Cases (15%)

Create multiple algorithms (IITPSx assembly) as test cases using the IITPSx instruction set (e.g., adding n numbers, computing a polynomial, mean, median, standard deviation, etc.). Include at least one test that exercises mul and div.

NOTE: the * symbol labels those tasks that are considered essential i.e. the minimum for a satisfactory standard.

.